

Program Structures and Algorithms

Spring 2023(SEC – 1)

NAME: Venkat Pavan Munaganti

NUID: 002722397

Task: Assignment 5: **Parallel Sorting**

Please see the presentation on *Assignment on Parallel Sorting* under the *Exams. etc.* module.

Your task is to implement a parallel sorting algorithm such that each partition of the array is sorted in parallel. You will consider two different schemes for deciding whether to sort in parallel.

1. A cutoff (defaults to, say, 1000) which you will update according to the first argument in the command line when running. It's your job to experiment and come up with a good value for this cutoff. If there are fewer elements to sort than the cutoff, then you should use the system sort instead.
2. Recursion depth or the number of available threads. Using this determination, you might decide on an ideal number (t) of separate threads (stick to powers of 2) and arrange for that number of partitions to be parallelized (by preventing recursion after the depth of $\lg t$ is reached).
3. An appropriate combination of these.

There is a *Main* class and the *ParSort* class in the *sort.par* package of the INFO6205 repository. The *Main* class can be used as is but the *ParSort* class needs to be implemented where you see "TODO..." [it turns out that these TODOs are already implemented].

Unless you have a good reason not to, you should just go along with the Java8-style future implementations provided for you in the class repository.

You must prepare a report that shows the results of your experiments and draws a conclusion (or more) about the efficacy of this method of parallelizing sort. Your experiments should involve sorting arrays of sufficient size for the parallel sort to make a difference. You should run with many different array sizes (they must be sufficiently large to make parallel sorting worthwhile, obviously) and different cutoff schemes.

Experimental Analysis:

Ran the experiment with different cut off values and threads to find out which combination gives the best runtime. Below are the inputs and different values used for the experiment

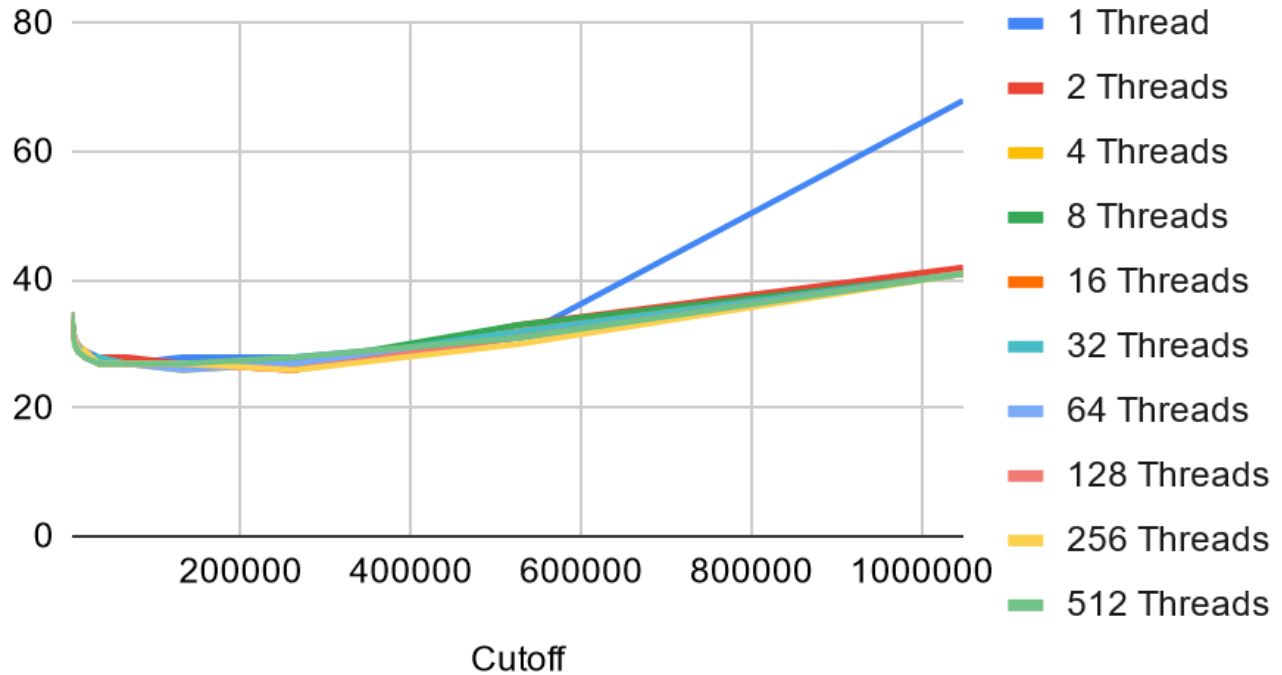
1. **Input size:** Ran the experiment with different yet sufficiently large input sizes ranging from 2^{20} to 2^{24}
2. **Cut off values:** Used different cut off values for each input size. Incrementing cut off values in decrementing fractions of powers of 2 to input size
3. **Threads:** Incrementing the threads in powers of 2 for each input size to find the optimal thread allocation for parallel sorting

Below are the results of the experiment,

For input size $N=1048576$,

| Parallel sort for input N=1048576 | | | | | | | | | | |
|-----------------------------------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|
| Cutoff | 1 T | 2 T | 4 T | 8 T | 16 T | 32 T | 64 T | 128 T | 256 T | 512 T |
| 1048576 | 68 | 42 | 41 | 41 | 41 | 41 | 41 | 41 | 41 | 41 |
| 524288 | 31 | 33 | 32 | 33 | 32 | 32 | 31 | 31 | 30 | 31 |
| 262144 | 28 | 26 | 27 | 27 | 26 | 27 | 27 | 26 | 26 | 28 |
| 131072 | 28 | 27 | 26 | 27 | 27 | 27 | 26 | 27 | 27 | 27 |
| 65536 | 27 | 28 | 27 | 27 | 27 | 27 | 27 | 27 | 27 | 27 |
| 32768 | 28 | 28 | 27 | 27 | 28 | 28 | 27 | 27 | 27 | 27 |
| 16384 | 29 | 28 | 29 | 28 | 28 | 28 | 28 | 28 | 29 | 28 |
| 8192 | 29 | 30 | 30 | 29 | 29 | 29 | 30 | 29 | 29 | 29 |
| 4096 | 31 | 31 | 30 | 31 | 31 | 31 | 31 | 31 | 31 | 30 |
| 2048 | 33 | 32 | 33 | 33 | 31 | 32 | 31 | 32 | 32 | 32 |
| 1024 | 33 | 35 | 34 | 33 | 34 | 35 | 33 | 34 | 34 | 35 |
| Average | 33.18181 | 30.90909 | 30.54545 | 30.54545 | 30.36363 | 30.63636 | 30.18181 | 30.27272 | 30.27272 | 30.45454 |

Runtime vs Cutoff for N=1048576




Analysis:

- Input size in this run is 2^{20}
- In above case, single thread is performing worse compared to multiple threads
- **Best cutoff:** Runtime is minimal when cutoff is almost equal to $\frac{1}{8}$ of input size, i.e, $N/8$
- **Best no of threads:** 64 Threads is giving the best runtime in parallel sort
- Since input size is comparatively less all threads above 4 are performing equally
- **Best combination:** When both methods are combined, threads in range 8 and 256, and cut off values in range 2^2 to 2^8 are giving overall best runtime

For input size $N=2097152$,

| Parallel sort for input N=2097152 | | | | | | | | | | |
|-----------------------------------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|
| Cutoff | 1 T | 2 T | 4 T | 8 T | 16 T | 32 T | 64 T | 128 T | 256 T | 512 T |
| 2097152 | 95 | 93 | 94 | 92 | 91 | 91 | 92 | 92 | 92 | 93 |
| 1048576 | 67 | 68 | 67 | 66 | 67 | 67 | 69 | 68 | 68 | 68 |
| 524288 | 57 | 56 | 57 | 56 | 56 | 57 | 57 | 56 | 60 | 58 |
| 262144 | 57 | 58 | 58 | 57 | 58 | 57 | 56 | 57 | 57 | 57 |
| 131072 | 57 | 56 | 56 | 57 | 56 | 57 | 56 | 56 | 56 | 56 |
| 65536 | 57 | 56 | 57 | 57 | 56 | 57 | 56 | 57 | 57 | 57 |
| 32768 | 59 | 59 | 58 | 59 | 58 | 59 | 56 | 58 | 57 | 58 |
| 16384 | 60 | 60 | 60 | 60 | 59 | 60 | 60 | 60 | 60 | 60 |
| 8192 | 62 | 62 | 64 | 62 | 62 | 62 | 62 | 63 | 62 | 62 |
| 4096 | 63 | 63 | 64 | 64 | 66 | 63 | 64 | 64 | 63 | 63 |
| 2048 | 66 | 66 | 67 | 67 | 66 | 67 | 66 | 64 | 67 | 67 |
| Average | 63.63636 | 63.36363 | 63.81818 | 63.36363 | 63.18181 | 63.36363 | 63.09090 | 63.18181 | 63.54545 | 63.54545 |




Analysis:

- Input size in this run is 2^{21}
- In above case, single thread is performing worse compared to multiple threads
- **Best cutoff:** Runtime is minimal when cutoff is almost equal to $\frac{1}{8}$ of input size, i.e, $N/8$
- **Best no of threads:** 64 Threads is giving the best runtime in parallel sort
- **Best combination:** When both methods are combined, threads in range 32 and 256, and cut off values in range $N/2^3$ to $N/2^6$ are giving overall best runtime

For input size $N=4194304$,

| Parallel sort for input N=4194304 | | | | | | | | | | |
|-----------------------------------|----------|----------|----------|----------|----------|----------|------|----------|----------|----------|
| Cutoff | 1 T | 2 T | 4 T | 8 T | 16 T | 32 T | 64 T | 128 T | 256 T | 512 T |
| 4194304 | 193 | 193 | 192 | 192 | 194 | 192 | 191 | 193 | 192 | 192 |
| 2097152 | 145 | 144 | 138 | 138 | 138 | 140 | 141 | 141 | 142 | 145 |
| 1048576 | 120 | 118 | 118 | 120 | 119 | 118 | 117 | 119 | 118 | 120 |
| 524288 | 118 | 117 | 117 | 119 | 118 | 118 | 117 | 118 | 118 | 118 |
| 262144 | 117 | 118 | 116 | 118 | 117 | 117 | 118 | 116 | 116 | 116 |
| 131072 | 120 | 117 | 120 | 117 | 120 | 120 | 119 | 117 | 117 | 119 |
| 65536 | 122 | 122 | 121 | 121 | 120 | 126 | 128 | 130 | 127 | 126 |
| 32768 | 135 | 138 | 131 | 131 | 130 | 137 | 137 | 131 | 133 | 130 |
| 16384 | 139 | 132 | 137 | 136 | 135 | 133 | 130 | 132 | 131 | 135 |
| 8192 | 143 | 140 | 144 | 140 | 146 | 142 | 135 | 139 | 148 | 140 |
| 4096 | 148 | 145 | 143 | 144 | 140 | 139 | 141 | 141 | 139 | 138 |
| Average | 136.3636 | 134.9090 | 134.2727 | 134.1818 | 134.2727 | 134.7272 | 134 | 134.2727 | 134.6363 | 134.4545 |

Runtime vs cutoff for N=4194304




Analysis:

- Input size in this run is 2^{22}
- In above case, single thread is performing worse compared to multiple threads
- **Best cutoff:** Runtime is minimal when cutoff is almost equal to $\frac{1}{8}$ of input size, i.e, $N/8$
- **Best no of threads:** 64 Threads is giving the best runtime in parallel sort
- **Best combination:** When both methods are combined, threads in range 16 and 256, and cut off values in range $N/2^2$ to $N/2^3$ are giving overall best runtime

For input size $N=8388608$,

| Parallel sort for input N=8388608 | | | | | | | | | | |
|-----------------------------------|----------|----------|----------|----------|----------|----------|----------|----------|-------|----------|
| Cutoff | 1 T | 2 T | 4 T | 8 T | 16 T | 32 T | 64 T | 128 T | 256 T | 512 T |
| 8388608 | 400 | 418 | 408 | 414 | 409 | 416 | 422 | 419 | 421 | 424 |
| 4194304 | 310 | 313 | 299 | 303 | 300 | 301 | 297 | 302 | 299 | 301 |
| 2097152 | 253 | 258 | 252 | 256 | 262 | 259 | 253 | 262 | 260 | 248 |
| 1048576 | 251 | 249 | 253 | 259 | 255 | 255 | 250 | 256 | 259 | 256 |
| 524288 | 259 | 260 | 256 | 249 | 243 | 245 | 253 | 246 | 247 | 244 |
| 262144 | 257 | 263 | 263 | 257 | 268 | 258 | 259 | 265 | 260 | 260 |
| 131072 | 278 | 269 | 269 | 272 | 270 | 265 | 275 | 272 | 279 | 276 |
| 65536 | 287 | 283 | 293 | 293 | 291 | 299 | 287 | 296 | 297 | 285 |
| 32768 | 297 | 305 | 306 | 299 | 310 | 308 | 310 | 311 | 302 | 298 |
| 16384 | 329 | 310 | 311 | 318 | 321 | 318 | 319 | 318 | 313 | 320 |
| 8192 | 334 | 324 | 329 | 327 | 333 | 329 | 326 | 326 | 341 | 327 |
| Average | 295.9090 | 295.6363 | 294.4545 | 295.1818 | 296.5454 | 295.7272 | 295.5454 | 297.5454 | 298 | 294.4545 |



Analysis:

- Input size in this run is 2^{23}
- In above case, single thread is performing worse compared to multiple threads
- **Best cutoff:** Runtime is minimal when cutoff is almost equal to $\frac{1}{8}$ of input size, i.e, $N/8$
- **Best no of threads:** 64 Threads is giving the best runtime in parallel sort
- **Best combination:** When both methods are combined, 64 threads, and cut off values $N/2^3$ is giving the best runtime

Conclusion:

- The experiment is performed in i7 11th generation, with 16 cores and 16th threads, when under turbo boost and full performance it could run efficiently till 64 threads considering the background tasks
- From all the above cases it is evident that runtime is minimal when cut off value is $\frac{1}{8}$ of the input size
- Runtime is minimal with **64 Threads**
- As the size of the array is increasing the runtime is minimal at cut off $\frac{1}{8}$ of the input size
- As the size of the array is increasing the runtime is minimal at cut off **64 Threads**

Main.java

```

import java.util.*;
import java.util.concurrent.Executor;
import java.util.concurrent.ForkJoinPool;

/**
 * This code has been fleshed out by Ziyao Qiao. Thanks very much.
 * CONSIDER tidy it up a bit.
 */
@xiaohuanlin
public class Main {
    @xiaohuanlin
    public static void main(String[] args) {
        processArgs(args);
        System.out.println("Degree of parallelism: " + ForkJoinPool.getCommonPoolParallelism());
        Random random = new Random();
        HashMap<Integer, List<List<Double>>> results = new HashMap<>();
        //Different array lengths in powers of 2
        for (int a = 20; a < 24; a++) {

            //Calculating the array size
            int size = (int) Math.pow(2, a);
            int[] array = new int[size];
            ArrayList<Long> timeList = new ArrayList<>();
            List<List<Double>> runtimes = new ArrayList<>();
            System.out.println("Parallel sort for input size N="+size);

            //Different cutoff values using doubling method, starting from 10000
            for (int j = 0; j <= 10; j++) {
                List<Double> runtime = new ArrayList<>();
                ParSort.cutoff = size/ (int) Math.pow(2, j);
                runtime.add((double) ParSort.cutoff);
                System.out.print(ParSort.cutoff+",");
                //Running on different threads in powers of 2 from 2^0 to 2^5 that is 1 to 32 threads
                for (int tp = 0; tp < 10; tp++) {
                    ForkJoinPool executor = new ForkJoinPool((int) Math.pow(2, tp));
                    long time;
                    long startTime = System.currentTimeMillis();

                    //10 runs to take average run time
                    for (int t = 0; t < 10; t++) {
                        //Filling the array with random values using random class

```

```

        for (int t = 0; t < 10; t++) {
            //Filling the array with random values using random class
            for (int i = 0; i < array.length; i++) array[i] = random.nextInt( bound: 10000000);
            ParSort.sort(array, from: 0, array.length);
        }
        long endTime = System.currentTimeMillis();
        time = (endTime - startTime);
        timeList.add(time);
        runtime.add((double) (time / 10));
        System.out.print((double) (time / 10)+",");
    }
    System.out.println();
    runtimes.add(runtime);
}
System.out.println();
System.out.println();
results.put(size, runtimes);
}

try {
    FileOutputStream fis = new FileOutputStream( name: "./src/parallelSort.txt");
    OutputStreamWriter isr = new OutputStreamWriter(fis);
    BufferedWriter bw = new BufferedWriter(isr);

    StringBuilder sb = new StringBuilder();
    for (Map.Entry<Integer, List<List<Double>>> e : results.entrySet()) {
        int size = e.getKey();
        List<List<Double>> runtimes = e.getValue();
        sb.append("parallel sort for input size N=" + size + "\n");
        for (List<Double> runtime : runtimes) {
            for(var i: runtime){
                sb.append(i+",");
            }
            sb.append("\n");
        }
        sb.append("\n\n");
    }
    bw.write(sb.toString());
    bw.flush();
    bw.close();
} catch (IOException e) {

```

```

        bw.close();
    } catch (IOException e) {
        e.printStackTrace();
    }
}

```

1 usage 👤 xiaohuanlin

```

private static void processArgs(String[] args) {
    String[] xs = args;
    while (xs.length > 0)
        if (xs[0].startsWith("-")) xs = processArg(xs);
}

```

1 usage 👤 xiaohuanlin

```

private static String[] processArg(String[] xs) {
    String[] result = new String[0];
    System.arraycopy(xs, srcPos: 2, result, destPos: 0, length: xs.length - 2);
    processCommand(xs[0], xs[1]);
    return result;
}

```

1 usage 👤 xiaohuanlin

```

private static void processCommand(String x, String y) {
    if (x.equalsIgnoreCase( anotherString: "N")) setConfig(x, Integer.parseInt(y));
    else
        // TODO sort this out
        if (x.equalsIgnoreCase( anotherString: "P")) //noinspection ResultOfMethodCallIgnored
            ForkJoinPool.getCommonPoolParallelism();
}

```

1 usage 👤 xiaohuanlin

```

private static void setConfig(String x, int i) {
    configuration.put(x, i);
}

```

/MismatchedQueryAndUpdateOfCollection/

```

private static final Map<String, Integer> configuration = new HashMap<>();

```

```

}

```

```
"C:\Program Files\Java\jdk-18.0.2.1\bin\java.exe" ...
```

Degree of parallelism: 15

Parallel sort for input size N=1048576

```
1048576,63.0,46.0,45.0,45.0,45.0,44.0,44.0,45.0,44.0,44.0,
524288,34.0,33.0,33.0,33.0,32.0,33.0,33.0,34.0,32.0,32.0,
262144,27.0,27.0,27.0,27.0,28.0,30.0,27.0,27.0,27.0,28.0,
131072,28.0,28.0,27.0,27.0,27.0,27.0,30.0,27.0,28.0,27.0,
65536,28.0,28.0,27.0,27.0,28.0,28.0,29.0,27.0,27.0,27.0,
32768,28.0,28.0,28.0,28.0,28.0,28.0,28.0,28.0,28.0,28.0,
16384,29.0,29.0,29.0,29.0,29.0,29.0,29.0,29.0,29.0,29.0,
8192,29.0,30.0,30.0,30.0,30.0,29.0,30.0,30.0,29.0,31.0,
4096,31.0,31.0,31.0,30.0,31.0,31.0,31.0,31.0,31.0,31.0,
2048,33.0,33.0,33.0,32.0,33.0,32.0,32.0,32.0,32.0,33.0,
1024,34.0,33.0,33.0,33.0,33.0,33.0,32.0,33.0,33.0,33.0,
```

Parallel sort for input size N=2097152

```
2097152,95.0,93.0,94.0,92.0,91.0,91.0,92.0,92.0,92.0,93.0,
1048576,67.0,68.0,67.0,66.0,67.0,67.0,69.0,68.0,68.0,68.0,
524288,57.0,56.0,57.0,56.0,56.0,57.0,57.0,56.0,60.0,58.0,
262144,57.0,58.0,58.0,57.0,58.0,57.0,57.0,57.0,57.0,57.0,
131072,57.0,56.0,56.0,57.0,56.0,57.0,56.0,56.0,56.0,56.0,
65536,57.0,56.0,57.0,57.0,56.0,57.0,56.0,57.0,57.0,57.0,
32768,59.0,59.0,58.0,59.0,58.0,59.0,58.0,58.0,57.0,58.0,
16384,60.0,60.0,60.0,60.0,59.0,60.0,60.0,60.0,60.0,60.0,
8192,62.0,62.0,64.0,62.0,62.0,62.0,62.0,62.0,62.0,62.0,
4096,63.0,63.0,64.0,64.0,65.0,63.0,64.0,64.0,63.0,63.0,
2048,66.0,66.0,67.0,67.0,66.0,67.0,66.0,64.0,67.0,67.0,
```

Parallel sort for input size N=4194304

```
4194304,193.0,193.0,192.0,192.0,194.0,192.0,191.0,193.0,192.0,192.0,
2097152,145.0,144.0,138.0,138.0,138.0,140.0,141.0,141.0,142.0,145.0,
1048576,120.0,118.0,118.0,120.0,119.0,118.0,117.0,119.0,118.0,120.0,
524288,118.0,117.0,117.0,119.0,118.0,118.0,118.0,118.0,118.0,118.0,
262144,117.0,118.0,116.0,118.0,117.0,117.0,117.0,116.0,116.0,116.0,
131072,120.0,117.0,120.0,117.0,120.0,120.0,119.0,117.0,117.0,119.0,
65536,122.0,122.0,121.0,121.0,120.0,126.0,128.0,130.0,127.0,126.0,
```

ParSort.java

```
// xiaohuanlin
public static void sort(int[] array, int from, int to) {
    if (to - from < cutoff) Arrays.sort(array, from, to);
    else {
        // FIXME next few lines should be removed from public repo.
        CompletableFuture<int[]> parsort1 = parsort(array, from, to: from + (to - from) / 2); // TO IMPLEMENT
        CompletableFuture<int[]> parsort2 = parsort(array, from: from + (to - from) / 2, to); // TO IMPLEMENT
        CompletableFuture<int[]> parsort = parsort1.thenCombine(parsort2, (xs1, xs2) -> {
            int[] result = new int[xs1.length + xs2.length];
            // TO IMPLEMENT
            int i = 0;
            int j = 0;
            for (int k = 0; k < result.length; k++) {
                if (i >= xs1.length) {
                    result[k] = xs2[j++];
                } else if (j >= xs2.length) {
                    result[k] = xs1[i++];
                } else if (xs2[j] < xs1[i]) {
                    result[k] = xs2[j++];
                } else {
                    result[k] = xs1[i++];
                }
            }
            return result;
        });
    }

    parsort.whenComplete((result, throwable) -> System.arraycopy(result, srcPos: 0, array, from, result.length));
    System.out.println("# threads: " + ForkJoinPool.commonPool().getRunningThreadCount());
    parsort.join();
}
}
```

2 usages xiaohuanlin

```
private static CompletableFuture<int[]> parsort(int[] array, int from, int to) {
    return CompletableFuture.supplyAsync(
        () -> {
            int[] result = new int[to - from];
            // TO IMPLEMENT
            System.arraycopy(array, from, result, destPos: 0, result.length);
            sort(result, from: 0, to: to - from);
            return result;
        }
    );
}
```